

# Global Illumination I




70001 – Advanced Computer Graphics: Photographic Image Synthesis

Abhijeet Ghosh

Lecture 14, Feb. 20<sup>th</sup> 2024


1

## Cornell Box – direct illumination



2

## Cornell Box – global illumination



3

## Rendering Equation [Kajiya 86]


- $$L_r(x, \omega_r) = L_e(x, \omega_r) + \int_{\Omega} f_r(x, \omega_r, \omega_i) L_i(x, \omega_i) \cos\theta_i d\omega_i.$$

- $L_e$  is the emitted radiance
- No explicit visibility term!

- Energy balance!


- $\Phi_r - \Phi_i = \Phi_e - \Phi_a$

- Path tracing for computing LTE
  - Monte Carlo integration



4


## Surface Form of LTE



- Integral over **area** instead of directions
  - $L(x' \rightarrow x) = L(x', \omega_r)$
  - $f_l(x'' \rightarrow x' \rightarrow x) = f_l(x', \omega_r, \omega_l)$

5


## Surface Form of LTE



- Geometric coupling term
  - $G(x' \leftrightarrow x) = V(x' \leftrightarrow x) | \cos \theta | | \cos \theta' | / \| x - x' \|^2$
- LTE as area integral
  - $L(x' \rightarrow x) = L_\theta(x' \rightarrow x) + \int_A f_l(x'' \rightarrow x' \rightarrow x) L(x'' \rightarrow x') G(x'' \leftrightarrow x') dA(x'')$

6

## Integrals over paths



- $L(x_1 \rightarrow x_0) = L_e(x_1 \rightarrow x_0) +$   
 $\int_A L_e(x_2 \rightarrow x_1) f_A(x_2 \rightarrow x_1 \rightarrow x_0) G(x_2 \leftrightarrow x_1) dA(x_2) +$   
 $\int_A \int_A L_e(x_3 \rightarrow x_2) f_A(x_3 \rightarrow x_2 \rightarrow x_1) G(x_3 \leftrightarrow x_2)$   
 $f_A(x_2 \rightarrow x_1 \rightarrow x_0) G(x_2 \leftrightarrow x_1) dA(x_3) dA(x_2) + \dots$

7

## Path tracing

- $L(x_1 \rightarrow x_0) = \sum_{i=0}^{\infty} P(x_i)$   
 $= P(x_1) + P(x_2) + \sum_{i=3}^{\infty} P(x_i)$ 
  - need to terminate path after finite terms!
  - need to sample paths for MC estimate of integral

8

## Path termination – Russian roulette

- Integrand not evaluated with probability  $q$ 
  - approx. with constant  $c$
- Evaluate integrand with probability  $1 - q$ , weighted by  $1/(1 - q)$ :

$$F' = \begin{cases} (F - qc)/(1 - q) & u > q \\ c & \text{otherwise} \end{cases}$$

- $E(F') = (1 - q) ((E(F) - qc)/(1 - q)) + qc$   
 $= E(F)$
- Same expected value but you pay the price of higher variance!

9

## Path termination

- Russian roulette allows us to terminate a path after finite terms
- $L(x_1 \rightarrow x_0) = P(x_1) + P(x_2) + 1/(1 - q) \sum_{i=3}^{\infty} P(x_i)$
- Alternatively:  
 $L(x_1 \rightarrow x_0) = 1/(1 - q_1) (P(x_1) + 1/(1 - q_2) (P(x_2) + 1/(1 - q_3) (P(x_3) + \dots$
- Path usually terminated from 3<sup>rd</sup> bounce onwards due to lower contribution
  - higher variance of low contribution is tolerable.

10

## Path sampling

- Sample light sources by **area**, and vertices according to **BRDF**
- Need to convert solid angle probability  $p_\omega$  to area  $p_A$

$$p_A = p_\omega \cdot (|\cos\theta_i| / \|x_i - x_{i-1}\|^2)$$

- MC estimate for path is then given as:

$$\frac{L_e(x_i \rightarrow x_{i-1})}{p_A(x_i)} \prod_{j=1}^{i-1} \frac{f_r(x_{j+1} \rightarrow x_j \rightarrow x_{j-1}) |\cos\theta_j|}{p_\omega(x_{j+1} - x_j)}$$

11

## Path tracing from camera



56 samples per pixel

12

## Bidirectional path tracing





Veach&Guibas 94

also see  
Lafortune&Willems93


25 samples per pixel

13


## Bidirectional path tracing




(a)  $s = 0, t = 3$



(b)  $s = 1, t = 2$



(c)  $s = 2, t = 1$



(d)  $s = 3, t = 0$

14

## Bidirectional path tracing

Paths combined with  
multiple importance sampling



15

## Bidirectional path tracing



40 samples per pixel

16



## Metropolis Light Transport – Veach&Guibas 97



250 mutations per pixel

17

## MLT

- Mutate an existing path to obtain a new path!
  - accept mutation based on energy **balance**
- Advantage over path tracing
  - path re-use
  - local exploration

18

## Metropolis-Hastings algorithm

```
x = x0
for i = 1 to n
  x' = mutate(x)
  a = accept(x, x')
  record(x, (1-a) * weight)
  record(x', a * weight)
  if (random() < a)
    x = x'
```

19

## Detailed balance

- Mutations proposed with transition probability  $T(x \rightarrow x')$
- Mutations accepted with acceptance probability  $a(x \rightarrow x')$
- Random walk in equilibrium requires:  
$$f(x) T(x \rightarrow x') a(x \rightarrow x') = f(x') T(x' \rightarrow x) a(x' \rightarrow x)$$

- $$a(x \rightarrow x') = \min \left( 1, \frac{f(x') T(x' \rightarrow x)}{f(x) T(x \rightarrow x')} \right)$$


20

## Mutation strategies

- Local perturbations
  - Good for explorations of high importance paths
- Proposals from a distribution
  - Good for maintaining **ergodicity**
  - perturbations can get stuck in a local minima!

21


## Lens perturbation



- Ideal for sub-paths of the form  $(L|D)DS^*E$  (context free grammar)

22


## Caustic perturbation



- Ideal for sub-paths of the form  $(D|L)S^*DE$  (context free grammar)

23


## Multi-chain perturbation




- Lens perturbation followed by caustic perturbation
  - Sampling caustics in a pool of water!

24


## MLT vs path tracing




Bidirectional PT



MLT



Path tracing



MLT

25



## Start-up bias

- MLT assumes sample  $x$  to already be drawn from the target distribution  $p$
- Need to draw many samples before achieving the target distribution
  - initial samples have to be discarded
  - wastage of samples ☹
- Draw from a proposal distribution  $q$  and weight the samples with  $f/q$ 
  - sub-optimal

26

## Energy Re-distribution Path Tracing – Cline et al. 05

- Combining MC path tracing and MLT mutations!
- Redistribution of energy from MC samples over image plane
  - lens perturbation!
- MLT has better convergence than MC PT for high dimensional problems
  - but worse for low dimensional problems!



27

## ERPT

---

```

EnergyRedistributionSampling ()
  for each integral domain,  $\Omega_i$ 
    for  $j = 1$  to  $m$ 
      create an MC sample,  $\bar{x}$ , in  $\Omega_i$  according to  $S_p$ 
      evaluate  $\mathbf{X}_f(\bar{x}) = f(\bar{x})/p(\bar{x})$ 
      if  $\mathbf{X}_f(\bar{x}) > 0$ 
        redistribute the energy of  $\mathbf{X}_f(\bar{x})$  using
          a balanced energy flow filter.
    
```

---

28


## Equal deposition flow

```
EqualDepositionFlow ( $\bar{x}$ ,  $e$ ,  $m$ ,  $e_d$ )  
  numChains =  $\lfloor \text{random}(0, 1) + e / (m \times e_d) \rfloor$   
  for  $i = 1$  to numChains  
     $\bar{y} = \bar{x}$   
    for  $j = 1$  to  $m$   
       $\bar{z} = \text{mutate}(\bar{y})$   
      if  $q(\bar{y} \rightarrow \bar{z}) \geq \text{random}(0, 1)$   
         $\bar{y} = \bar{z}$   
      deposit  $e_d$  energy at  $\bar{y}$ 
```

- Here mutation implies selecting another pixel in a local neighborhood and depositing some of the sample's energy to its neighbor
  - Image pixels in a neighborhood will all converge to similar energy

29


## Path tracing



100 samples per pixel

30


## MLT



100 mutations per pixel

31

## ERPT





36 MC samples +  
50 mutations per pixel

32



## Correlated visibility sampling – [Ghosh & Heidrich 06]

- Similar in spirit to ERPT
  - but applied to visibility sampling for direct illumination
- Bidirectional sampling for un-occluded regions – 1<sup>st</sup> pass
  - Mark partially occluded pixels
- Metropolis energy exchange on the image plane for partially occluded regions – 2<sup>nd</sup> pass



33

## Correlated visibility sampling



Bidirectional sampling  
20 SIR samples

Visibility sampling  
16 SIR + 16 MCMC samples

34

## Filtering MC Noise in Rendering

[Kalantari et al. 2015]




- Neural network is trained to predict the appropriate parameters of a filter kernel which is then applied to de-noise a noisy MC rendering output.
- Training uses additional features of the image besides pixel colors such as surface normal, surface positions, texture as features for predicting the kernel

35

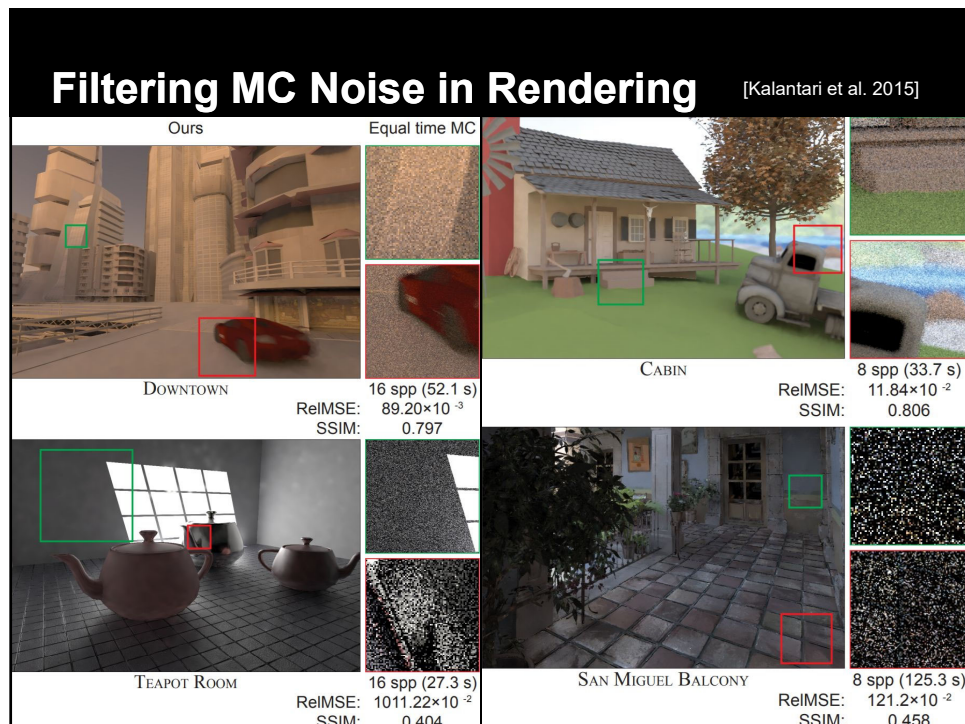
## Filtering MC Noise in Rendering

[Kalantari et al. 2015]



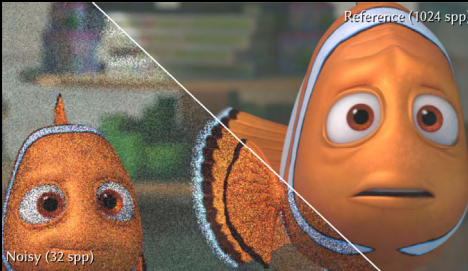
- Neural network is trained to predict the appropriate parameters of a filter kernel which is then applied to de-noise a noisy MC rendering output.
- Training uses additional features of the image besides pixel colors such as surface normal, surface positions, texture as features for predicting the kernel

36

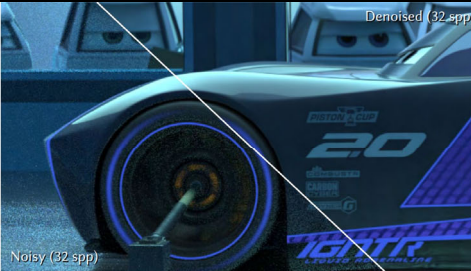


37

## Filtering MC Noise with CNNs [Bako et al. 2017]



TRAINING



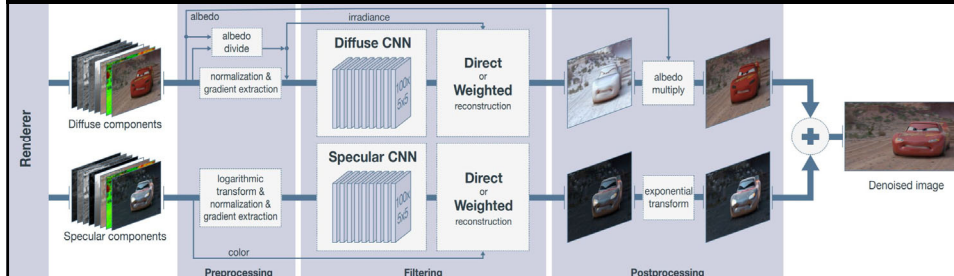
TEST

- CNN trained on noisy and converged frame pairs from Finding Dory
- Trained CNN can be applied to other films with different style of lighting and appearance such as Cars 3.
- Denoising stable enough to be used in production rendering pipelines to create final quality animation frames.

38

# Filtering MC Noise with CNNs

[Bako et al. 2017]



- Separate CNN networks trained for denoising diffuse and specular renderings
- Diffuse shading is denoised after factoring out the albedo, which is post-multiplied.

39

# Filtering MC Noise with CNNs

[Bako et al. 2017]


Ours	Input (32 spp)		
relative $\ell_2$ 1 - SSIM	19.21e-3 0.354	relative $\ell_2$ 1 - SSIM	14.92e-3 0.360
relative $\ell_2$ 1 - SSIM	18.88e-3 0.271	relative $\ell_2$ 1 - SSIM	20.31e-4 0.069
		<p>The CNN trained on Finding Dory generalizing to other types of scenes in animation film production</p>	
relative $\ell_2$ 1 - SSIM	9.28e-3 0.090		

40



## Filtering MC Noise with CNNs

[Bako et al. 2017]



CNN training is a bit content specific. Hence re-training is required for different type of rendered input, e.g., indoor scenes as opposed to Pixar animations. After re-training on the appropriate type of content, the method generalized well.