

DEPARTMENT OF COMPUTING

IMPERIAL COLLEGE OF SCIENCE, TECHNOLOGY AND MEDICINE

---

# Popular Network Architectures (& BatchNorm)

---

*LeNet-5 [4], MNIST, AlexNet [3], ImageNet, VGG [5], Inception (GoogleLeNet), BatchNorm [2],  
ResNet [1], DenseNet, Squeeze-Excite Net U-Net, Data Augmentation*

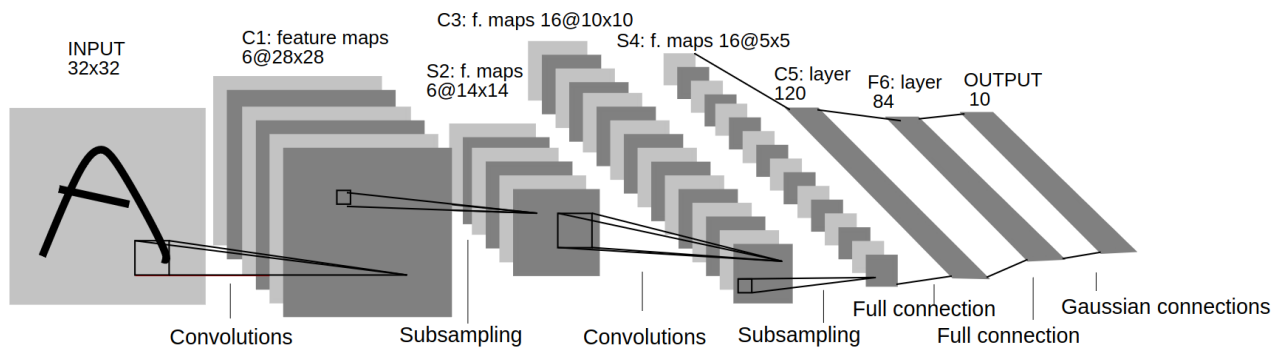
*Author: Anton Zhitomirsky*

## Contents

**1 LeNet-5**

**2**

# 1 LeNet-5



**Figure 1:** Architecture of LeNet-5 a Convolutional Neural Network, here for digits recognition. Each plane is a feature map, i.e. a set of units whose weights are constrained to be identical [4]

```

1 # from https://pytorch.org/tutorials/beginner/blitz/neural_networks_tutorial.html
2 import torch
3 import torch.nn as nn
4 import torch.nn.functional as F
5
6
7 class Net(nn.Module):
8
9     def __init__(self):
10         super(Net, self).__init__()
11         # 1 input image channel, 6 output channels, 5x5 square convolution
12         # kernel
13         self.conv1 = nn.Conv2d(1, 6, 5)
14         self.pool1 = nn.MaxPool2d(2)
15         self.conv2 = nn.Conv2d(6, 16, 5)
16         self.pool2 = nn.MaxPool2d(2)
17         # an affine operation: y = Wx + b
18         self.fc1 = nn.Linear(16 * 5 * 5, 120) # 5*5 from image dimension
19         self.fc2 = nn.Linear(120, 84)
20         self.fc3 = nn.Linear(84, 10)
21
22     def forward(self, x):
23         x = self.conv1(x)
24         x = F.relu(x)
25         x = self.pool1(x)
26         x = self.conv2(x)
27         x = F.relu(x)
28         x = self.pool2(x)
29
30         x = x.torch.flatten(x,1) # flatten all dimensions except the batch dimension
31
32         x = self.fc1(x)
33         x = F.relu(x)
34         x = self.fc2(x)
35         x = F.relu(x)
36         x = self.fc3(x)
37
38         return x

```

code/LetNet.py

LeNet [4] was initially designed for low-resolution, black and white image recognition, specifically for digits. It demonstrated that CNNs could reliably perform both tasks of object localization and recognition (on low-resolution black and white images).

The last steps of the LeNet-5 architecture employ fully connected layers to convert features into final predictions. The scenario of MNIST data didn't matter, because of its small number of output classes. This complicates the architecture when scaling up, influencing both computational resources and architecture.

## References

- [1] Kaiming He et al. *Deep Residual Learning for Image Recognition*. 2015. arXiv: 1512.03385 [cs.CV]. URL: <https://arxiv.org/abs/1512.03385>.
- [2] Sergey Ioffe and Christian Szegedy. *Batch Normalization: Accelerating Deep Network Training by Reducing Internal Covariate Shift*. 2015. arXiv: 1502.03167 [cs.LG]. URL: <https://arxiv.org/abs/1502.03167>.
- [3] Alex Krizhevsky, Ilya Sutskever, and Geoffrey E. Hinton. "ImageNet classification with deep convolutional neural networks". In: *Commun. ACM* 60.6 (May 2017), pp. 84–90. ISSN: 0001-0782. DOI: 10.1145/3065386. URL: <https://papers.nips.cc/paper/2012/file/c399862d3b9d6b76c8436e924aPaper.pdf>.
- [4] Y. Lecun et al. "Gradient-based learning applied to document recognition". In: *Proceedings of the IEEE* 86.11 (1998), pp. 2278–2324. DOI: 10.1109/5.726791. URL: [http://vision.stanford.edu/cs598\\_spring07/papers/Lecun98.pdf](http://vision.stanford.edu/cs598_spring07/papers/Lecun98.pdf).
- [5] Karen Simonyan and Andrew Zisserman. *Very Deep Convolutional Networks for Large-Scale Image Recognition*. 2015. arXiv: 1409.1556 [cs.CV]. URL: <https://arxiv.org/pdf/1409.1556.pdf>.