
Self-supervised Learning

This week's paper is [\[1\]](#)

Author: Anton Zhitomirsky

Contents

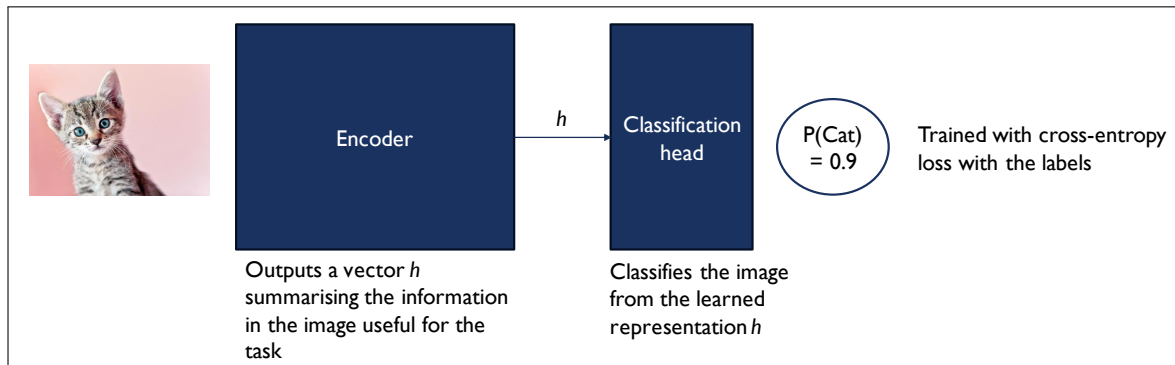
1	Motivation	3
1.1	Standard supervised classification pipeline	3
1.2	Self-supervised learning classification pipeline	3
1.2.1	Proxy	3
2	Contrastive-based Learning — SIMCLR	4
2.1	Main Idea	4
2.2	Main Components	5
2.3	Walkthrough	5
2.4	Augmentation Pipeline	6
2.5	Loss Function — NT-XNet loss	7
2.5.1	Similarity Measure	7
2.5.2	Loss	7
2.5.3	Triplet Loss	8
2.6	Evaluation of Self-supervised Learning	8
2.6.1	Finetuning	8
2.6.2	Linear evaluation	9
2.7	BatchSize	9
2.8	Follow-up work	10
2.8.1	Bootstrap your won latent	10
2.8.2	DINO	11
2.9	Conclusion	11
3	Generative-based Learning	11
3.1	Main-Idea	11
3.1.1	Vision transformers (refresher)	12
3.2	Masked Auto-encoders are scalable vision learners	12
3.2.1	Loss	13
3.2.2	Scalability	13

3.3 Conclusion	13
4 Joint-Embedding Prediction	13
4.1 i-Jepa	13
4.1.1 Loss	15
5 Examples in Healthcare	15
Bibliography	16

1 Motivation

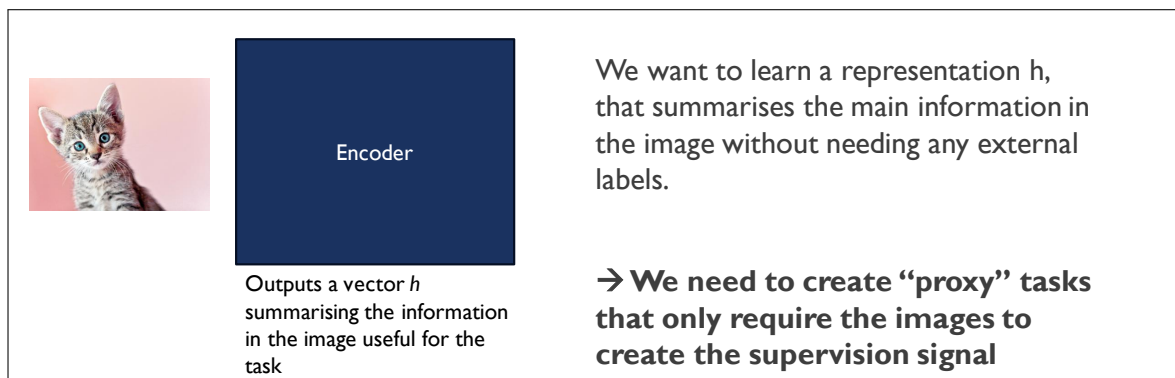
Supervised classifiers are very data hungry. This is a problem where labels are very expensive to acquire, but we have a lot of unlabelled images available (e.g. healthcare). The question is, could we learn to understand images without any labels?

1.1 Standard supervised classification pipeline



The question remains, ‘can we also learn useful image representation h without labels?’

1.2 Self-supervised learning classification pipeline



We chop off the classification head because we don't have a label, so there is no classification task, yet we still want to output a vector which summarises an image which still has all the important information to be used for classification as a downstream task. We require an objective for the network to strive for. Therefore we create proxies

1.2.1 Proxy


How can we create synthetic tasks that will get the network to learn useful representation?

- **Contrastive-based learning:** teaching the network to recognise meaningful pairs of images (Section 2).
- **Generative-based learning:** teaching the network to reconstruct an image from a corrupted version (Section 3)
- **Joint-Embedding Prediction:** a mix of both (Section 4).


2 Contrastive-based Learning — SIMCLR

2.1 Main Idea

A meaningful representation should put similar images closer to each other.



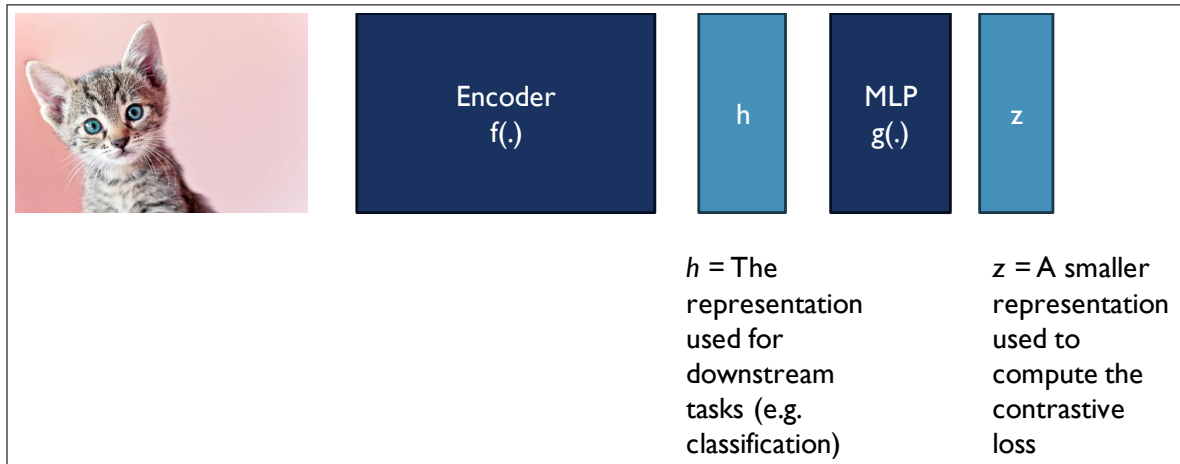
However, without labels I don't know which images are dogs or cats so I cannot use this information as supervision signal!



The extra cat here is a blurred version of a cat already assigned to the circular space. Therefore, this perturbed version of the same image should have a similar representation to the original image compared to all other images.

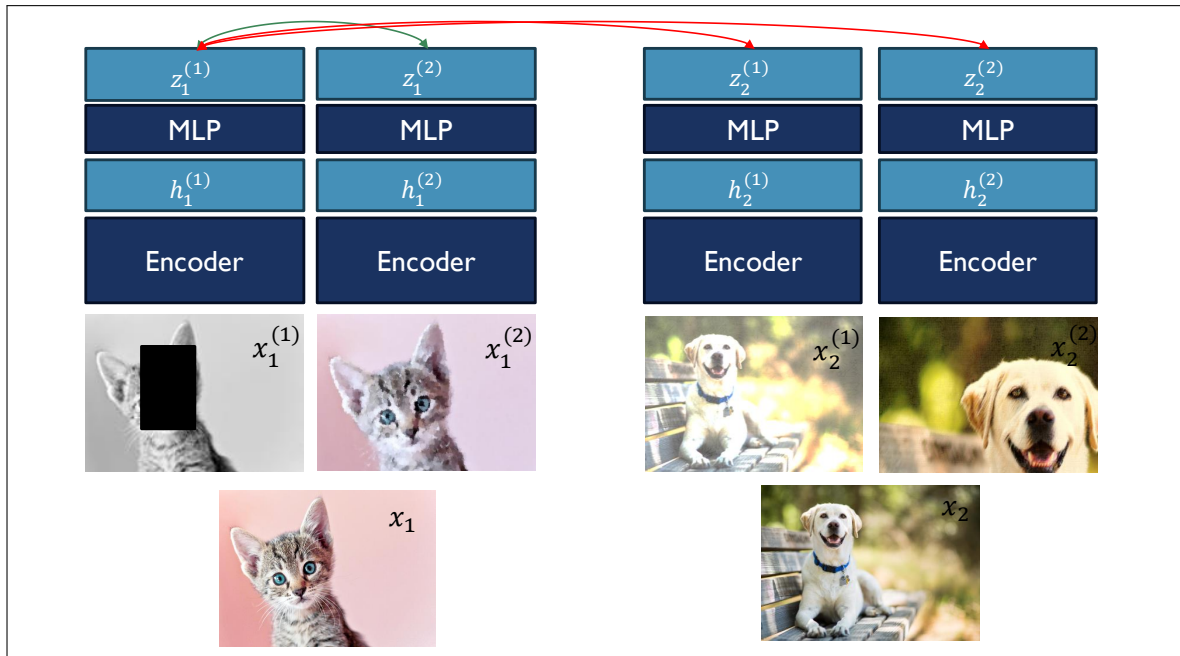
This is the main idea of contrastive-based learning: create multiple versions of the same image and teach the network to recognise the correct pair. The paper is at [\[1\]](#).

2.2 Main Components

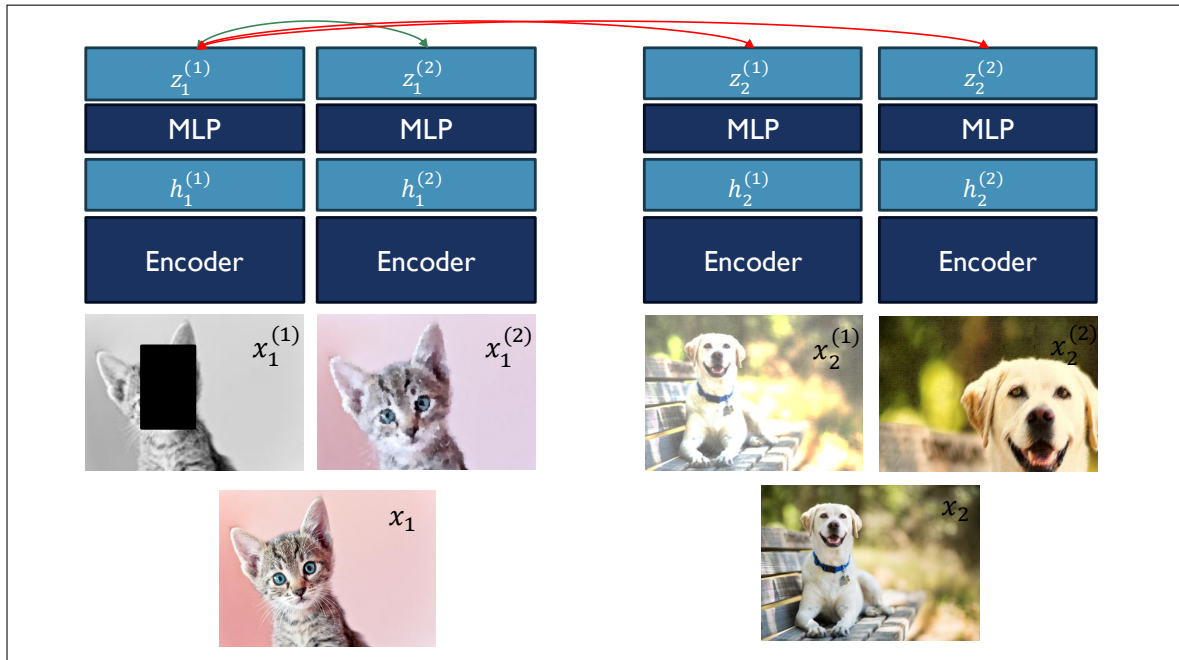


1. We are trying to pretrain the encoder $f(\cdot)$ that we are trying to train
2. This encoder outputs a representation h
3. Add small multi-layer perception to make the dimensions smaller. This is because in high dimensional space, comparing similarities is harder, and not necessarily robust. This projector translates from vectors of size 2048 to 128 for example.

2.3 Walkthrough



Start with two images $x_1 \wedge x_2$ and apply augmentations to them to produce (for example) two views. Then pass these images through the network. We want the embeddings of the two views of the cat to be close to each other, and at the same time have them far away from the dogs.



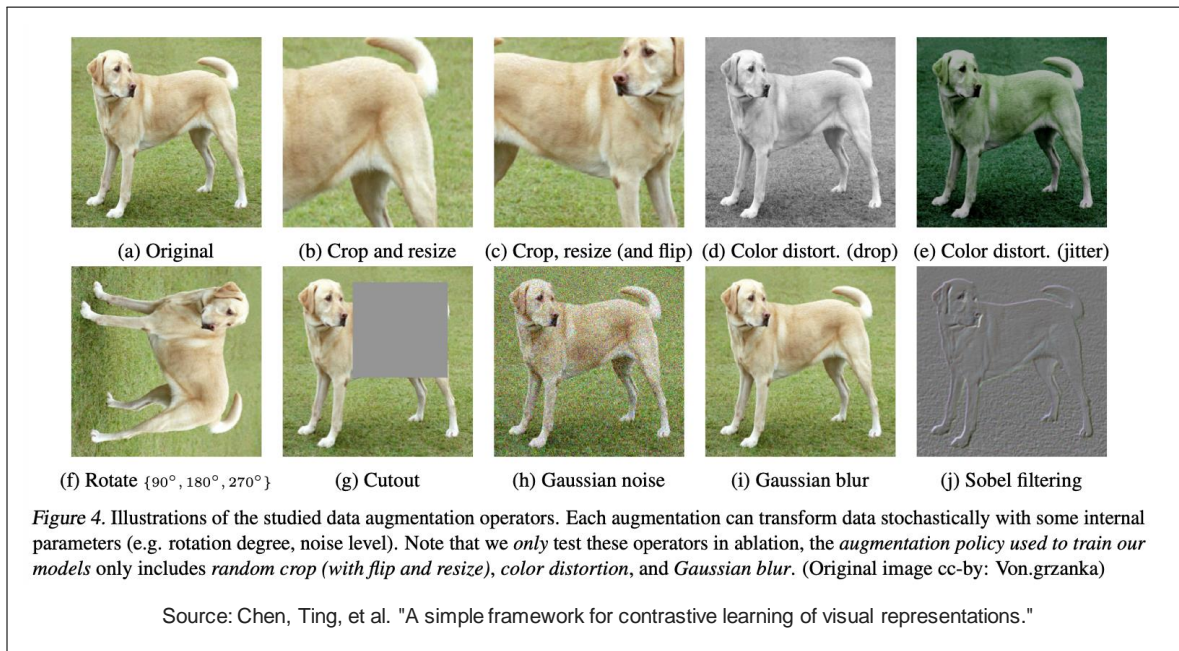
This is done for both of the views

2.4 Augmentation Pipeline

The augmentation you choose has very big consequences on what your model will actually learn. For example, an augmentation that removes colour will implicitly teach the network that you should not care about the colour, only about the concept (shapes and textures).

We require the augmentation pipeline to:

- Needs to reflect what information the model should disregard and what it should focus on
- Needs to be hard enough, otherwise trivial information is learnt.



These augmentations are very heavy, yet in the paper it is shown that this is what makes the model work

2.5 Loss Function — NT-XNet loss

2.5.1 Similarity Measure

The problem they try to solve is ‘how can we attract positive pairs while repelling negative pairs’. This is the **normalised temperature contrastive loss** algorithm.

$$\text{sim}(\mathbf{u}, \mathbf{v}) = \frac{\mathbf{u}^\top \mathbf{v}}{\|\mathbf{u}\| \cdot \|\mathbf{v}\|} \quad (2.5.1)$$

This is essentially the cosine similarity which is the angle between the two vectors.

2.5.2 Loss

For each positive pair (i, j) we have the following loss:

$$\ell_{i,j} = -\log \frac{\exp(\text{sim}(\mathbf{z}_i, \mathbf{z}_j)/\tau)}{\sum_{k=1}^{SN} \mathbf{1}_{[k \neq i]} \exp(\text{sim}(\mathbf{z}_i, \mathbf{z}_k)/\tau)} \quad (2.5.2)$$

1. We want to minimise the loss, therefore we want to maximise the logarithm (since there is a leading – in the loss function)
2. Therefore, we need to maximise the numerator and minimise the denominator.
3. **numerator:** We want to maximise the similarity (attract the embeddings we want together). I.e. maximise the nominator.
4. **denominator:** Measure similarity between the first image and all other images in the batch. I.e. minimise the denominator.
5. τ is the temperature, which controls how much to penalise hard negatives (negative pairs wrongly mapped close to each other). A low temperature penalises them more.

The final loss is computed across all positive pairs, both (i, j) and (j, i) , in a mini-batch, by averaging all $\ell_{i,j}$.

Concretely to build the loss for the entire batch with N samples one needs to follow the following process:

1. **Compute all projected embeddings for the first view** of each sample $\mathbf{z}^{(1)} = [z_1^{(1)}; z_2^{(1)}; \dots; z_N^{(1)}]$, vector of size $[N, \text{feature_dim}]$. For that first apply augmentation, pass through encoder E, finally pass through MLP projector to get z.
2. **Repeat it for the second view** $\mathbf{z}^{(2)} = [z_1^{(2)}; z_2^{(2)}; \dots; z_N^{(2)}]$
3. Then **gather all representations in one big vector** (to compute the similarities for the denominator), of size $[2N, \text{feature_dim}]$

$$\mathbf{z} = [\mathbf{z}^{(1)}, \mathbf{z}^{(2)}] = [z_1^{(1)}; \dots; z_N^{(1)}; z_1^{(2)}; \dots; z_N^{(2)}]$$

4. In this vector the positive pairs will be elements i and $i + N$ (and the opposite) Hence the **final loss averaged over all positive pairs** will be:

$$L_{\text{Batch}} = \frac{1}{2N} \sum_{i=1}^N [l_{i,i+N} + l_{i+N,i}]$$

2.5.3 Triplet Loss

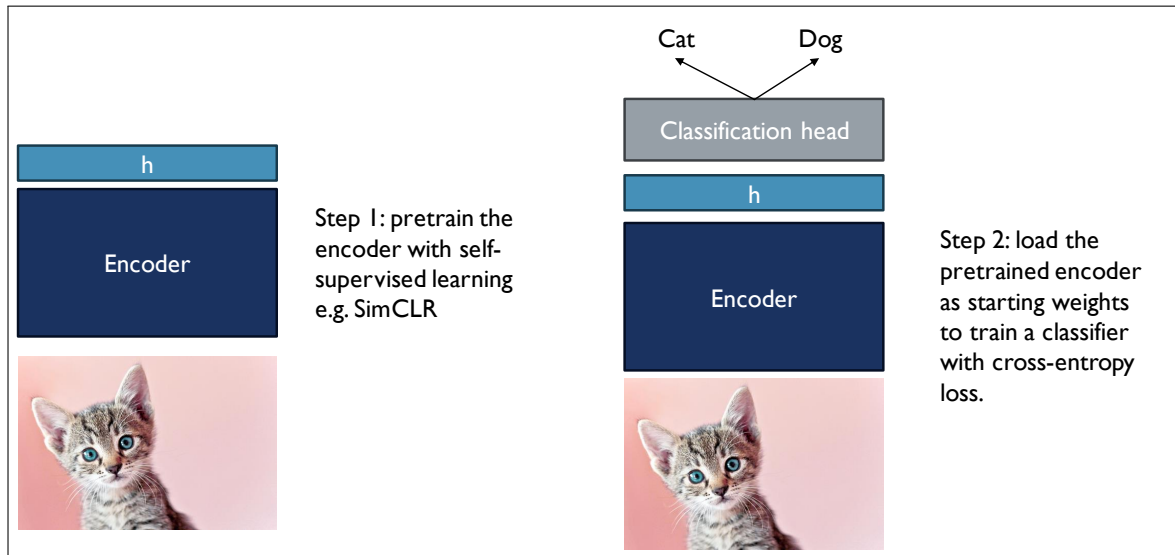
$$\mathcal{L}_{triplet}(x, x^+, x^-) = \sum_{x \in \mathcal{X}} \max(0, \|f(x) - f(x^+)\|_2^2 - \|f(x) - f(x^-)\|_2^2 + \epsilon) \quad (2.5.3)$$

- x is one image
- x^+ is the corresponding positive
- x^- is a randomly selected negative image
- ϵ is the margin parameter controlling how far the negatives should be compared to the positive pairs (temperature τ plays this role in NT-Xent loss)

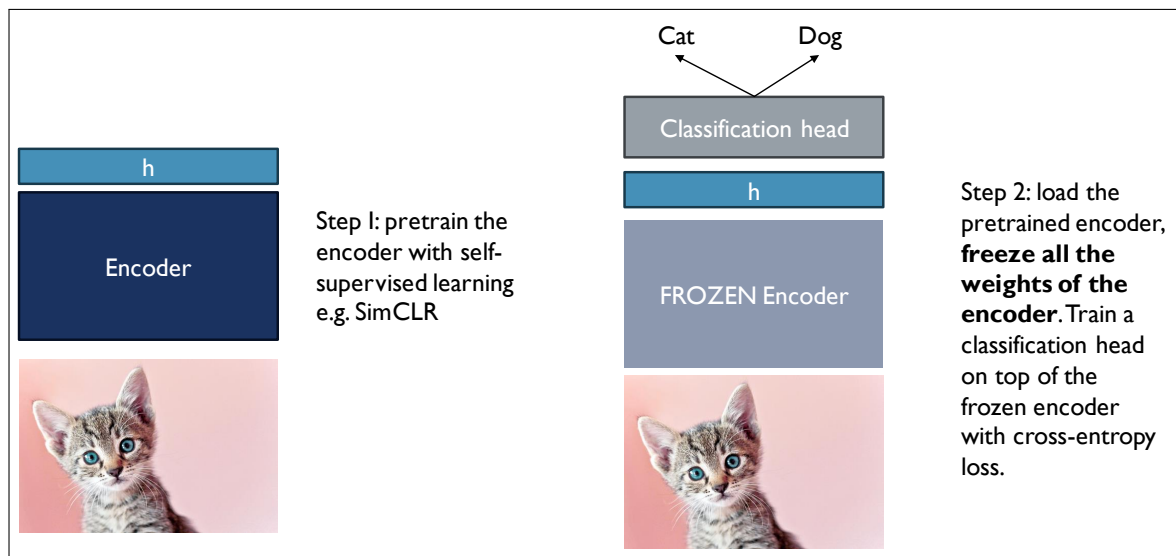
Here you want a similar thing: you want the positive pairs to be closer than the distance between the negatives (plus the margin). This quantity is negative so the loss is zero. “The positive pairs should be closer to each other by at least ϵ margin compared to the distance to any negative pair”

2.6 Evaluation of Self-supervised Learning

2.6.1 Finetuning



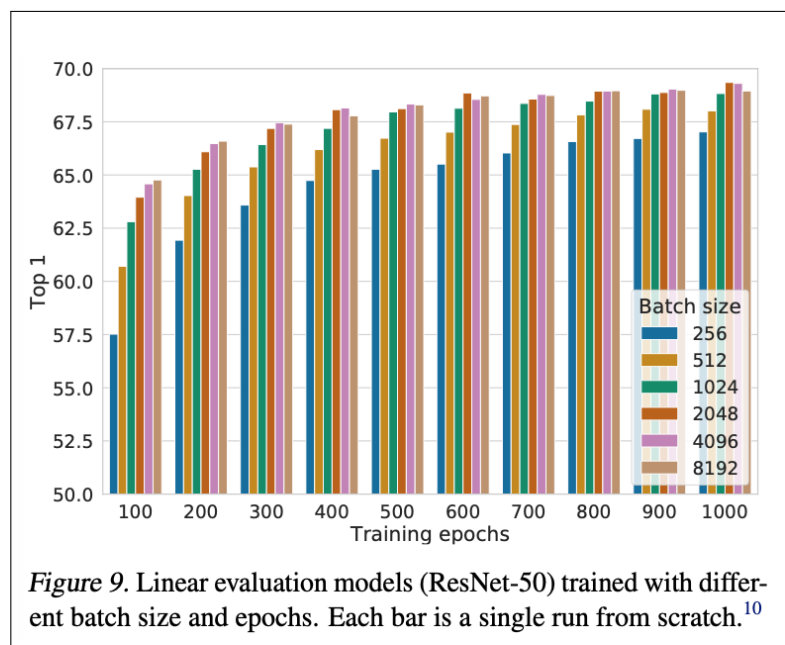
2.6.2 Linear evaluation



“you can only base yourself on what you learned during the pre-training. So it can be a good way to check what you’re actually learning during the pretraining as opposed to if you do for fine tuning you might modify the encoder” If you have few labels, then sometimes it’s better to keep the frozen encoder and just retain the classification label layer

2.7 BatchSize

To learn good representation, it is crucial to use very large batch sizes with SimCLR.




Linear evaluation — Key take-away; the more you pre-train the model, the better it gets. Also, bigger batch sizes learn a better representation of the image, the more images there are in your batch, the harder the task is.

2.8 Follow-up work

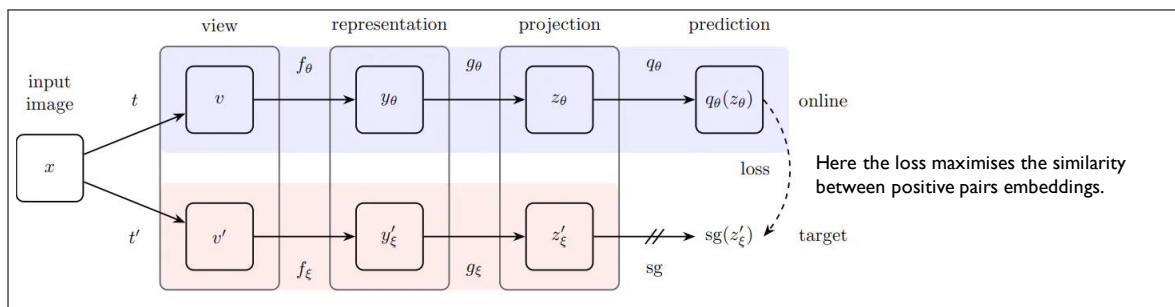
SimCLR requires very large batch size to provide good representation, this means very high computational requirements (expensive, difficult to train).

2.8.1 Bootstrap your won latent


BYOL removes the need to use negative pairs in the loss function, instead one just optimises the similarity on the positive pairs \rightarrow less need for bigger batch sizes. For that they need to introduce a new architecture to avoid learning trivial embeddings (otherwise all converge to the same point).



Where previously you would pass your two images into the same path, here we have them go into two separate paths (because they might obtain the trivial solution). The student and teacher (as described above) the weights will not be the same so they won't return the same representation




"To further avoid the risk of still learning the trivial solution where these two just end up converging to the same point, they add yet another MLP (prediction) and try to match the output of online and target. You're trying to say that the output of the student network after the prediction should be the same as the output of the target network without the prediction"



2.8.2 DINO

Very similar to training BYOL, but using vision transformers as encoders. “Because they use vision transformers they have access to this attention layer and so they can visualise the attention map of the network when they create the self-supervised encoding, and by visualising these attention maps they can generate self-supervised segmentation maps”.



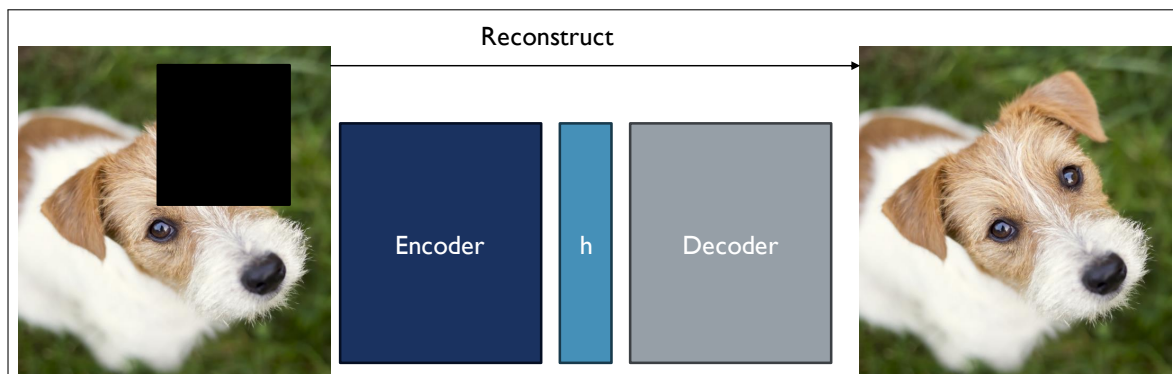
2.9 Conclusion

Contrastive learning is one pretext task that has proven very useful in self-supervised learning, but it's not the only way. In particular, it has some drawbacks such as large batch sizes, design of the augmentation pipeline etc.

3 Generative-based Learning

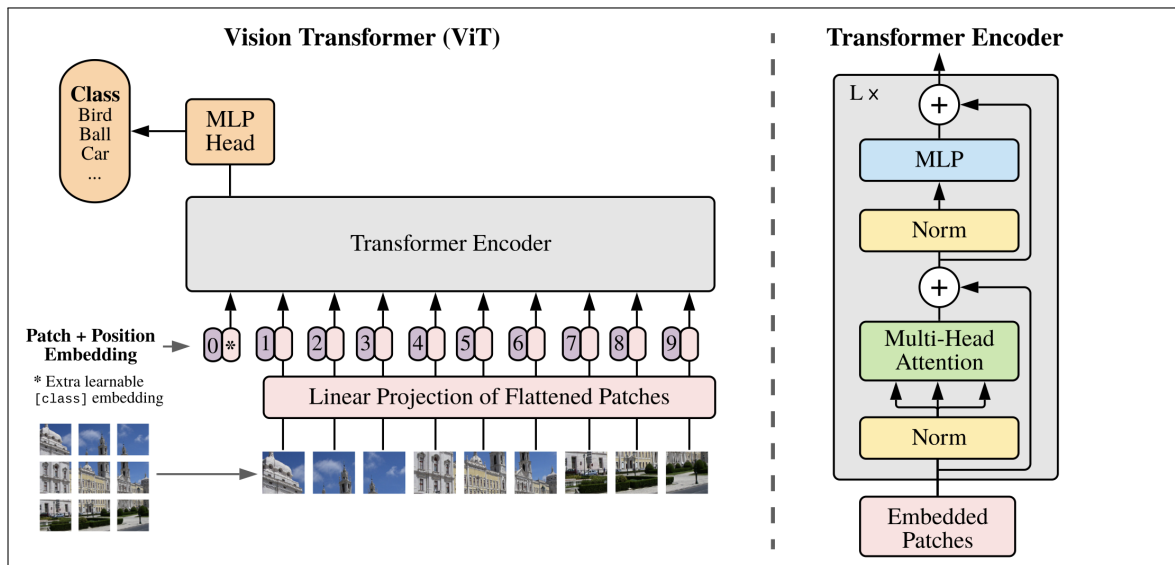
3.1 Main-Idea

Another way to use the image as supervision signal is to teach the network to reconstruct the original image from a masked image.



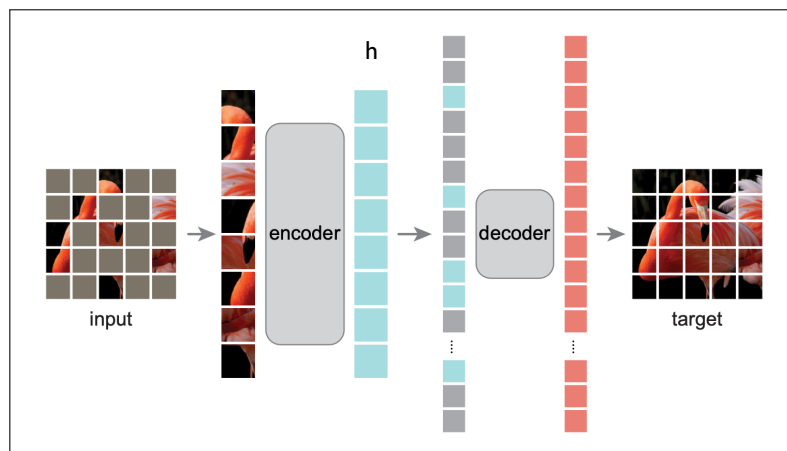
Instead of using pairs, constructing pairs to teach the network to learn semantics, they teach the network to reconstruct the original image from a masked image.

3.1.1 Vision transformers (refresher)




“They start with one image, they cut it in different non-overlapping patches then they pass it through some embedding and then they get it into a transformer encoder, which has some attention layers, and the attention layers update the representation of that patch using the information from the other patches, and progressively it gets updated and then it learns useful representation of each of these single patches.”

3.2 Masked Auto-encoders are scalable vision learners



They divide the image in multiple subpatches and randomly mask some of the patches. After taking the patches they have, they feed it through the encoder which is a vision transformer. The encoder gets a representation for each patch and then you fill in the blanks with some mask token you learn during the process. The decoder then predicts each patch.

The encoder here is drawn bigger because most of the computation is done in the encoder, and the decoder is pretty small. The goal is to put most of the details in the encoder (after training the decoder is thrown away).

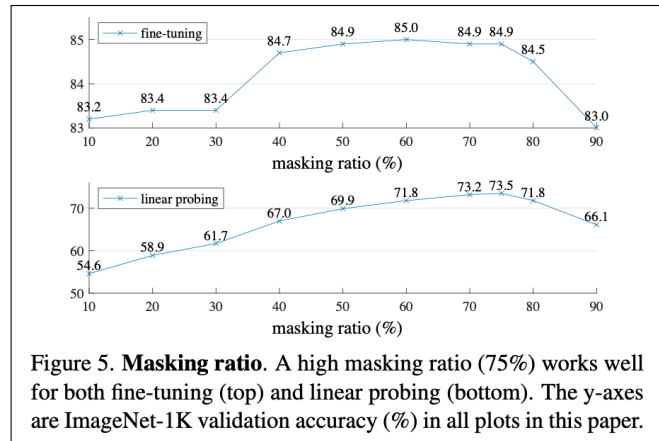


The reconstruction of the detail is pretty bad, but this is because the size of the decoder is smaller (and we throw away the decoder after training)

3.2.1 Loss

$$MSE = \sum (\hat{x}_i - x_i)^2, \quad \text{where } i \text{ is the pixel index} \quad (3.2.1)$$

3.2.2 Scalability



The best result they got was approximately 75% mask ratio

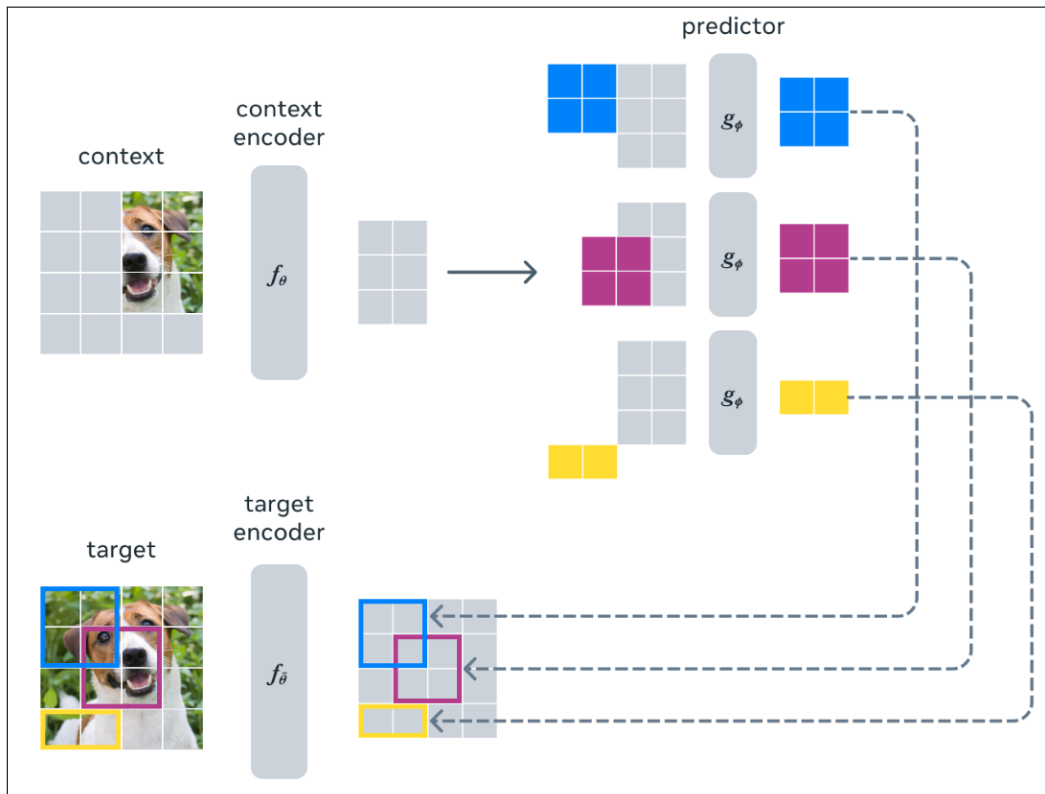
3.3 Conclusion

MAE avoids some of the drawbacks of CL, however, training a full reconstruction model is quite expensive. Maybe its not necessarily to full ylearn a perfect decoder model since we throw it away after training?

4 Joint-Embedding Prediction

4.1 i-Jepa

Idea is to combine strength of CL and MAE approaches.



Some patches are masked, and we pass on the available patches to the context encoder. At the same time, we sample a region of the blue window, and pass it through the target embedding. Repeat this for several regions. The goal of this paradigm is that instead of predicting the missing image region, we are going to predict the encoded representation because the goal of the encoders is to summarise information, so instead of wasting compute reconstructing every single pixel in the image, we are going to reconstruct the important information which there should be if the network is trained summarised.

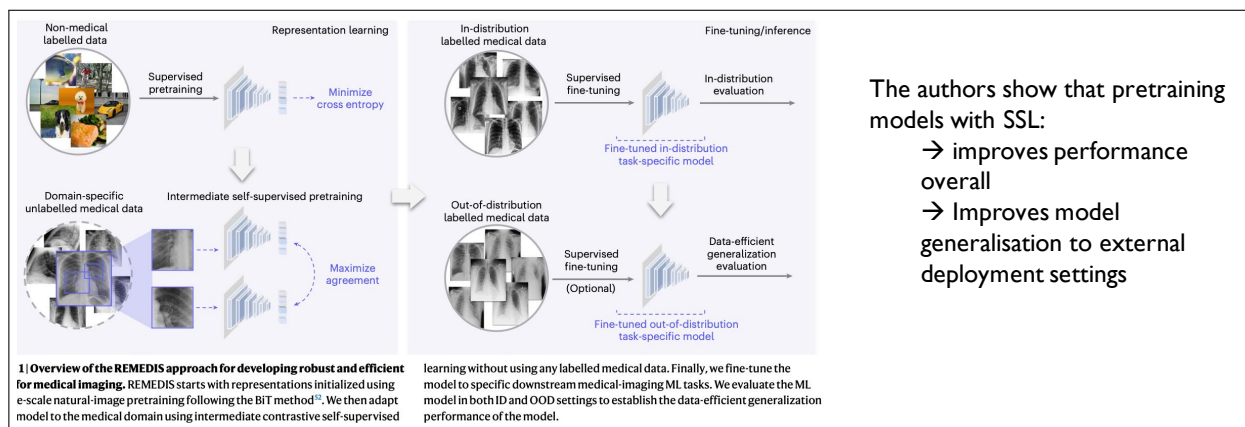


4.1.1 Loss

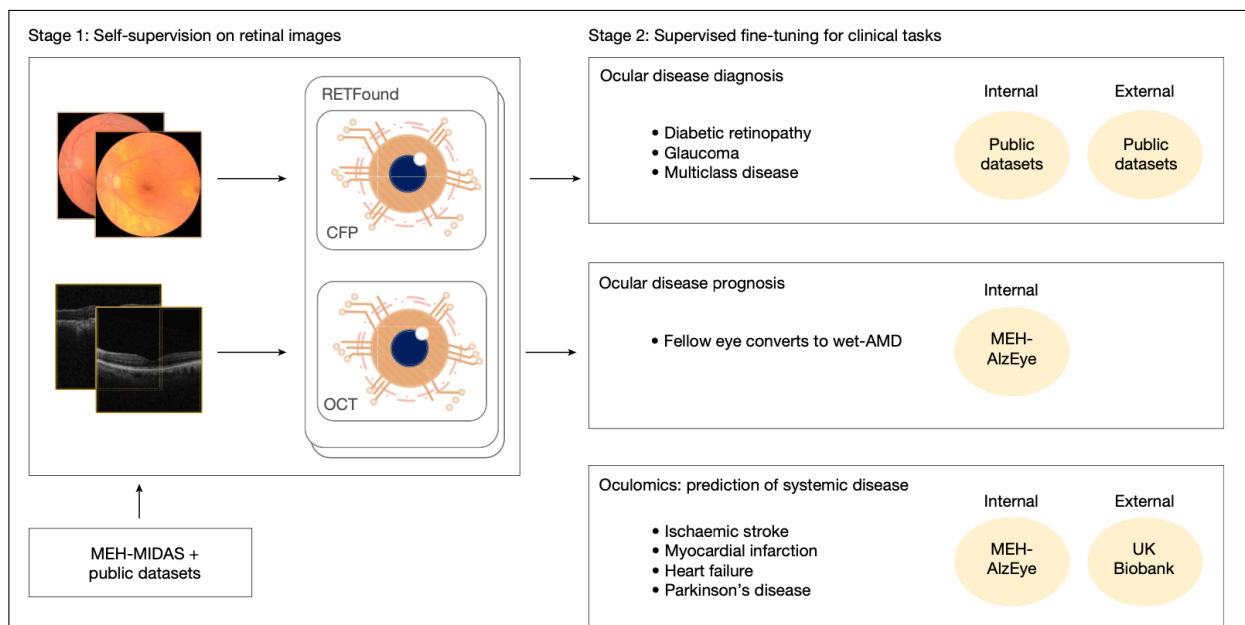
The loss is simply the average L_2 distance between the predicted patch-level representations $\hat{s}_y(i)$ and the target patch-level representation $s_y(i)$ i.e.,

$$\frac{1}{M} \sum_{i=1}^N D(\hat{s}_y(i), s_y(i)) = \frac{1}{M} \sum_{i=1}^M \sum_{j \in B_i} \|\hat{s}_{y_j} - s_{y_j}\|_2^2. \quad (4.1.1)$$

5 Examples in Healthcare



1. “Start by pretraining it on ImageNet (better performance than from scratch)”
2. “then they run SIMCLR from these weights. The goal is to go from a very generic encoder that has no idea about medical languages to specialise it to learn on these unlabelled chest x-rays”
3. “then they will use this on downstream tasks”



Half a million scans were trained which didn't have labels, and saw that pre-training improved performance than from scratch.

References

- [1] Ting Chen et al. *A Simple Framework for Contrastive Learning of Visual Representations*. 2020. arXiv: [2002.05709](#) [[cs.LG](#)].