

# example

---

*example description*

*Author: Anton Zhitomirsky*

## Contents

<b>1</b>	<b>NLP Classification tasks</b>	<b>2</b>
1.1	Natural Language Inference . . . . .	2
<b>2</b>	<b>Naive Bayes</b>	<b>2</b>
2.1	Add-one smoothing . . . . .	3
2.2	Binary Naive Bayes . . . . .	4
2.3	Controlling for negation . . . . .	4
2.4	Problems . . . . .	4
2.5	Summary . . . . .	4
<b>3</b>	<b>Logistic Regression</b>	<b>4</b>
3.1	Multiple Classes . . . . .	5
3.2	Summary . . . . .	5
<b>4</b>	<b>Neural Networks (NNs)</b>	<b>5</b>
4.1	Docuemnt Representation . . . . .	5
4.2	Neural Networks . . . . .	6
<b>5</b>	<b>Recurrent neural networks (RNNs)</b>	<b>6</b>
5.1	Vanishing gradient problem . . . . .	6
<b>6</b>	<b>CNNs</b>	<b>6</b>
<b>7</b>	<b>Accuracy and F1</b>	<b>7</b>
7.1	Macro averaging . . . . .	7
7.2	Micro averaging . . . . .	7

# 1 NLP Classification tasks

**Definition 1.1** (Classification).

$$\hat{y} = \arg \max_y P(y|x) \quad (1)$$

Predicting which ‘class’ an observation belongs to

A Model produces a score (logit), a sigmoid makes this between 0 and 1, and 0.5 is our decision boundary. In multi-class classification we then use softmax.

## 1.1 Natural Language Inference

a model is presented with a pair of sentences and must classify the relationship between their meanings. For example in the MultiNLI corpus, pairs of sentences are given one of 3 labels: entails, contradicts and neutral. These labels describe a relationship between the meaning of the first sentence (the premise) and the meaning of the second sentence (the hypothesis). Here are representative examples of each class from the corpus:

- **Premise:** The kitten is climbing the curtains again
- **Hypothesis:** The kitten is sleeping
- **Entailment:** If the hypothesis is implied by the premise
- **Contradiction:** If the hypothesis contradicts the premise (in this example, the hypothesis is contradicting)
- **Neutral:** otherwise (neither is necessarily true)

# 2 Naive Bayes

(Generative Algorithm)

**Definition 2.1** (Bayes Rule).

$$\underbrace{P(y|x)}_{\text{Posterior}} = \frac{\overbrace{P(x|y)}^{\text{Likelihood}} \overbrace{P(y)}^{\text{Prior}}}{\underbrace{P(x)}_{\text{Evidence}}}$$

Since  $P(x)$  won't change for different classes

👉 why?

$$\hat{y} = \arg \max_y P(y|x) = \arg \max_y P(x|y)P(y)$$

We can further state that since  $x$  is a set of features  $x_1, \dots, x_I$  we have a independence assumption:

**Definition 2.2** (Naive Bayes Classifier).

$$P(x_1, \dots, x_I|y) = P(x_1|y) \cdots P(x_I|y)$$

$$\hat{y} = \arg \max_y P(x_1, \dots, x_I|y)P(y) = \arg \max_y P(y) \prod_{i=1}^I P(x_i|y)$$

## 2.1 Add-one smoothing

Raw input is transformed into a numerical representation - i.e. each input  $x$  is represented by a feature vector by using a Bag of Words approach (count of each word in the input) “This was another good movie for holiday watchers. There was a nice little twist at the end” Collect statistics from our training

	a	about	another	and	...	was	you
Review #1	1	0	1	0		2	0

data (find what  $P(y|x)$  is) after performing limited data-preprocessing.

Training corpus	good	movie	bad	class
another <b>good</b> <b>movie</b> for holiday watchers . a little twist from the ordinary scrooge movie . enjoyable .	1	1	0	+
it seems like just about everybody has made a christmas carol movie . others are just <b>bad</b> and the time period seems to be perfect .	0	0	1	+
if you 're looking for the same feel <b>good</b> one but in a new setting , this one 's for you .	1	0	0	+
this is a first for me , i didn 't like this <b>movie</b> . it was really <b>bad</b> .	0	1	1	-
it was <b>good</b> but the christmas carol by dickens was emotionally moving .	1	0	0	-

**Figure 1:** Example of training corpus. Here, we are only concerned with the words ‘good’, ‘movie’, and ‘bad’ for classes ‘+’ and ‘-’

$$P(y) \rightarrow P(+) = \frac{3}{5}, \quad P(-) = \frac{2}{5}$$

$$P(\text{good}|+) = \frac{2}{4}$$

i.e.  $\frac{\text{frequency of word for class}}{\text{total count of words for this class}}$

this introduces the problem that one of our probabilities could be zero; therefore, adjust naive bayes classifier with ‘Add-one smoothing’

**Definition 2.3** (Add-one smoothing).

$$P(x_i|y) = \frac{\text{count}(x_i, y) + 1}{\sum_{x \in V} (\text{count}(x, y) + 1)} = \frac{\text{count}(x_i, y) + 1}{(\sum_{x \in V} \text{count}(x, y)) + |V|}$$

$$P(\text{good}|+) = \frac{2+1}{4+3} = \frac{3}{7} \quad P(\text{good}|-) = \frac{1+1}{3+3} = \frac{2}{6}$$

$$P(\text{movie}|+) = \frac{1+1}{4+3} = \frac{2}{7} \quad P(\text{movie}|-) = \frac{1+1}{3+3} = \frac{2}{6}$$

$$P(\text{bad}|+) = \frac{1+1}{4+3} = \frac{2}{7} \quad P(\text{bad}|-) = \frac{1+1}{3+3} = \frac{2}{6}$$

Therefore for a test example: “Not as **good** as the old **movie**, rather **bad**” we use Equation 2.2 to get  $P(+ )P(x|+) = \frac{3}{5} \times \frac{3 \times 2 \times 2}{7^3} = 0.021$  and  $P(-)P(x|-) = \frac{2}{5} \times \frac{2 \times 2 \times 2}{6^3} = 0.014$ . So final outcome of the model is +. However, this sentence is clearly negative, yet we classify as positive.

## 2.2 Binary Naive Bayes

Within binary naive bayes, we only consider if a feature is present, rather than considering every time it occurs. Note, this means re-calculating the conditional probabilities from the training data.

E.g. “Not as **good** as the old **movie**, rather **bad movie**” we have  $P(+ )P(x|+) = \frac{3}{5} \times \frac{3 \times 2 \times 2}{7^3} = 0.021$  and  $P(-)P(x|-) = \frac{2}{5} \times \frac{2 \times 2 \times 2}{6^3} = 0.014$ . If we were still operating under the same formula in Equation 2.2 then we would have iterated over the label ‘movie’ twice, and thus multiplied by  $P(good|+)$  twice (in the positive case).

## 2.3 Controlling for negation

We append ‘Not\_’ after any logical negation (e..g n’t, not, no, never) until the next punctuation mark.

“I didn’t like the movie, but it was better than Top Gun” becomes “I didn’t NOT\_like NOT\_the NOT\_movie, but it was better than Top Gun”

## 2.4 Problems

- Conditional independence assumption
- Features considered equally important
- Context of words not taken into account
- New words (not seen at training) cannot be used

## 2.5 Summary

Very quick to train, Some evidence it works well on very small datasets.

# 3 Logistic Regression

(Discriminative Algorithm) because we directly learn  $P(Y|X)$  since we don’t care about  $P(Y)$  or  $P(X)$ . They learn the input feauters most useful to discriminate between the different classes, without considering the likelihood of the input itself.

$$y(x) = g(z) = \frac{1}{1 + e^{-z}}$$

$$z = w \cdot x + b$$

$$P(y = 1) = \frac{1}{1 + e^{-(w \cdot x + b)}}$$

where  $w$  is how important an input feature is to the classification decision and we have a threshold at 0.5 for decision making.

Test example: Not as good as the old **movie**, rather **bad**.

$$x = [1.0, 1.0] \quad w = [-5.0, 2.5] \quad b = 0.1$$

$$\begin{aligned} P(y=1|x) &= g(z) \\ &= g([-5.0, 2.5] \cdot [1.0, 1.0] + 0.1) \\ &= g(-2.4) \\ &= 0.08 \end{aligned}$$

$$\begin{aligned} P(y=0|x) &= 1 - g(z) \\ &= 0.92 \end{aligned}$$



51

If we don't have the  $w$ , then we learn parameters to make the model predictions close to our labels by using a loss function (to measure the distance between the true and predicted labels) and optimization algorithm (to minimize the function usually through gradient descent.)

**Definition 3.1** (Logistic Regression).

$$H(P, Q) = - \sum_i P(y_i) \log Q(y_i)$$

### 3.1 Multiple Classes

Weights and bias are learnt per class, then use a softmax function over the result of finding  $z_i$ .

$$y = g(z_i) = \frac{e^{z_i}}{\sum_{j=1}^k e^{z_j}}$$

### 3.2 Summary

Logistic Regression considers the importance of features, so is better (than Naive Bayes) at dealing with correlated features, also better (than Naive Bayes) with larger datasets.

## 4 Neural Networks (NNs)

**Definition 4.1** (Linear Layer).

$$z = w \cdot x + b = \sum_{i=0}^I w_i x_i + b$$

**Definition 4.2** (Non-linear activation function).

$$y = g(z)$$

**Definition 4.3** (Fully-connected layers).

$$FFN(x) = (g^2(g^1(xW^1 + b^1))W^2 + b^2)W^3 + b^3$$

The neural networks allow for automatically learning dense feature representations, not one-hot encodings.

### 4.1 Document Representation

How to get a document representation of sentence of fixed dimensionality?

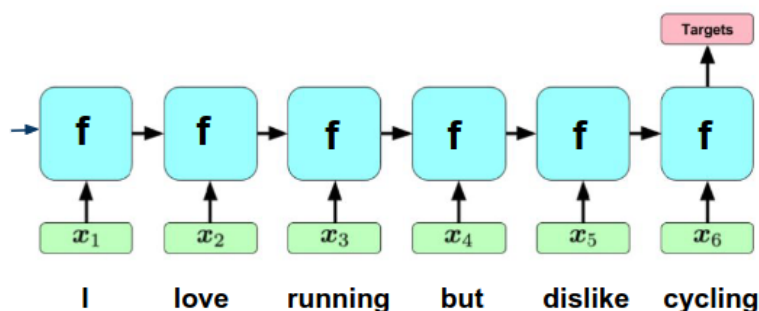
Average of sentence - bad idea:

Model architecture fixed to sentence length size, model weights learnt for specific word positions

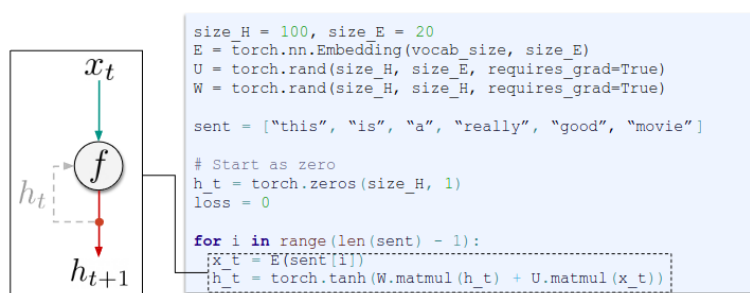
## 4.2 Neural Networks

Automatically learned features; flexibility to fit highly complex relationships in data, but they require more data to learn more complex patterns.

## 5 Recurrent neural networks (RNNs)



Natural language data is made up of sequences, so its natural to represent in a RNN; the value of a unit depends on own previous outputs - the last hidden state is the input to the output layer.



### 5.1 Vanishing gradient problem

The model is less able to learn from earlier inputs (Tanh derivatives are between 0 and 1) (Sigmoid derivatives are between 0 and 0.25) - Gradient for earlier layers involves repeated multiplication of the same matrix  $W$  - depending on the dominant eigenvalue this can cause gradients to either 'vanish' or 'explode'

👉 RNNs perform better when you need to understand longer range dependencies

## 6 CNNs

CNNs are composed of a series of convolution layers, pooling layers and fully connected layers. Convolutional layers Detect important patterns in the inputs. Pooling layers Reduce dimensionality of features and transform them into a fixed-size. Fully connected layers Train weights of learned representation for a specific task.

We can stack multiple filters on-top of eachother also.

👉 CNNs can perform well if the task involves key phrase recognition.

## 7 Accuracy and F1

**Definition 7.1** (accuracy).

$$Accuracy = \frac{TP + TN}{TP + FP + TN + FN}$$

**Definition 7.2** (f1-measure).

$$f1 = 2 \times \frac{precision \times recall}{precision + recall} = \frac{TP}{TP + 0.5(FP + FN)}$$

### 7.1 Macro averaging

averaging of each class F1 scores: increases the emphasis on less frequent classes

### 7.2 Micro averaging

TPs, TNs, FNs. FPs are summed across each class.

Microaveraged F1

		Predicted		
		Airplane	Boat	Car
Actual	Airplane	2	1	0
	Boat	0	1	0
	Car	1	2	3

		Predicted		
		Airplane	Boat	Car
Actual	Airplane	2	1	0
	Boat	0	1	0
	Car	1	2	3

$$\frac{\sum_i^C TP_i}{\sum_i^C TP_i + 0.5(\sum_i^C FP_i + \sum_i^C FN_i)} = \frac{\sum_i^C TP_i}{|Dataset|} = Accuracy$$